

# TOTAL RECALL

A Deep Dive into  
Ruby Memory  
Management





ruby memory issues



All

15,000,000 results

Settings

Tools

About

### Hunting Down Memory Issues In Ruby: A Definitive Guide | Toptal

<https://www.toptal.com> » Blog ▾

Memory issues in programs can be quite frustrating whether you're the user or the developer attempting to solve the problem. In Ruby, the garbage collector ...

### Debugging memory leaks in Ruby - Sam Saffron

<https://samsaffron.com/archive/2015/03/31/debugging-memory-leaks-in-ruby> ▾

Mar 30, 2015 - Debugging memory leaks in Ruby. At some point in the life of every Rails developer you are bound to hit a memory leak. It may be tiny amount of constant memory growth, or a spurt of growth that hits you on the job queue when certain jobs run.

### How to debug Ruby memory issues | RightScale Engineering Blog

<eng.rightscale.com/2015/09/16/how-to-debug-ruby-memory-issues.html> ▾

Sep 16, 2015 - Here at RightScale we have noticed recently issues with some of our server array boxes running dangerously low on memory. This can be ...

### Debugging a Memory Leak on Heroku | via @codeship

<https://blog.codeship.com/debugging-a-memory-leak-on-heroku/> ▾

Jun 19, 2017 - Reading Time: 10 minutes. This is one of the most frequent questions I'm asked by Heroku Ruby customers: "How do I debug a memory leak?".

### How I spent two weeks hunting a memory leak in Ruby - Fire Bowl

<www.be9.io/2015/09/21/memory-leak/> ▾

Sep 21, 2015 - A blog devoted to Ruby, Rails, Clojure, and other awesome tech.

### Finding the cause of a memory leak in Ruby - Stack Overflow

<https://stackoverflow.com/questions/.../finding-the-cause-of-a-memory-leak-in-ruby> ▾

4 answers

Dec 16, 2013 - It looks like you are entering The Lost World here. I don't think the problem is with c-bindings in racc either. Ruby memory management is both ...

Memory Leaks in my Ruby on Rails Application

3 answers Apr 26, 2016

How do I track down a memory leak in my Ruby code?

4 answers Jan 6, 2014

ruby on rails - How can I find a memory leak on Heroku?

4 answers Nov 12, 2012

ruby/ruby on rails memory leak detection

7 answers Oct 2, 2008

More results from stackoverflow.com

### What I Learned About Hunting Memory Leaks in Ruby 2.1

<blog.skylight.io/hunting-for-leaks-in-ruby/> ▾

Oct 30, 2014 - We'd been noticing a slow but persistent leak in our Rails app for a ... The most useful

WE GOT  
BIG  
REPUTATIONS



# A Little About Me



Matt Duszynski



@bsdzunk



[github.com/dzunk](https://github.com/dzunk)

- Software engineer at **weedmaps**<sup>®</sup>
- From Houston, TX
- Ruby developer since 2012
- Interests
  - Coffee
  - Music
  - Golf
- Docker contributor
  - OK, I updated the docs one time

**weedmaps**<sup>®</sup>




# What I'm Talking About

- Terms and concepts
- Under the hood
  - Page allocation
  - Object allocation
  - Garbage collection
  - Visualizing the heap
- Improvements, past and future
- Troubleshooting memory consumption
  - Leak or bloat?
  - Using memory efficiently
  - Tools for the toolbox



# What I'm **NOT** Talking About




**TRUFFLE  
RUBY**

- Other interpreters have different memory models
- The concepts will overlap (JVM has garbage collection too)
- Specific terms, variables, libraries, etc. will not apply across platforms

# Let's Get On The Same Page

- Allocation
  - OS assigns memory to a Ruby process (`malloc`)
- Deallocation
  - Ruby process returns unused memory to the OS (`free`)
- Heap
  - The entire memory space allocated to a Ruby process
- Page
  - 16 kilobyte section of the heap
  - Smallest amount of memory that can be allocated
- Slot
  - 40 byte section of an individual page
  - One object per slot



# How it Works: Allocation

The layer cake of memory abstraction

- Ruby VM - `Object.new` in `ObjectSpace`
- Allocator - `malloc`
- Operating system - virtual memory
- Hardware - memory management unit

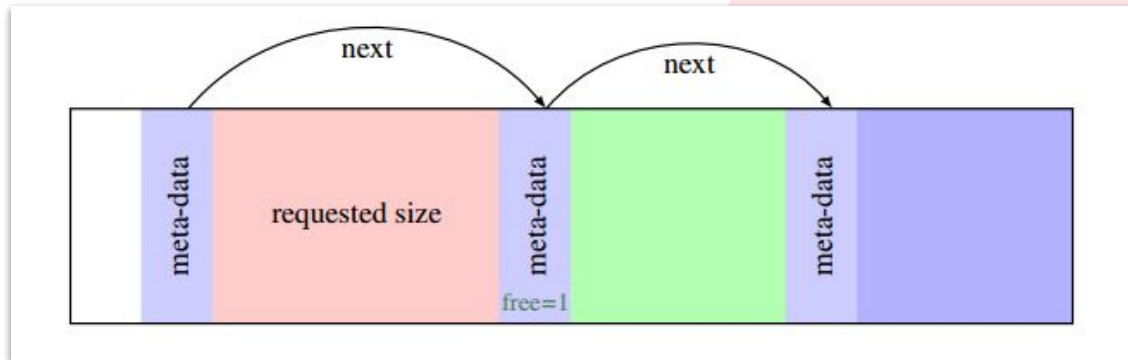
We're only talking about the first two



# How it Works: Page Allocation

When your process needs memory for anything:

1. Is there enough contiguous free space in the heap?
2. Run garbage collection
3. Is there enough contiguous free space in the heap?
4. Allocate additional pages





# How it Works: Object Allocation

ObjectSpace – Every living object in a Ruby process

```
[1] pry(main)> Object.new  
=> #<Object:0x00007ff5eaa76010>
```

Object address to binary:

111111111110101111010101010011101	10000000010000
Page address	Object ID

ObjectSpace entries are 40 bytes or less.  
Any non-trivial amount of data lives  
elsewhere in memory.



# How it Works: Object Allocation

`String.new "Lorem ipsum dolor sit amet..."` (445 chars)

1. Create an object in a slot in `ObjectSpace`
  - No room? Add a page with `malloc`
2. Allocate memory for the string itself
  - `malloc(445)`

The `ObjectSpace` entry is a pointer to the memory location from step two.


It's in the heap somewhere, but locations aren't guaranteed.

Causes memory fragmentation.

`ObjectSpace` does **not** represent all of Ruby's memory!

# How it Works: Garbage Collection

- Mark and sweep
  - Start from root objects
  - Mark all reachable (referenced) objects
  - Free all unmarked objects
  - Unmark all remaining objects
- Advantages
  - Transparent to running program
  - Handles cyclic references
- Disadvantages
  - Program must be suspended during GC runs
  - Leads to memory fragmentation over time without compaction



# How it Works: Generational GC

Most objects are dereferenced very soon after creation.

Mark and sweep only objects that have been created since the last GC run.


Faster and more efficient, but you eventually have to traverse all objects.

In Ruby, this is “minor GC” and “major GC”



# How it Works: Incremental GC

- Separates “mark” phase and “sweep” phase
- Distributes GC pause time more evenly
  - Better to pause 5 times for 3ms each than one time for 15ms
- Doesn't decrease overall pause time




# How it Works: Visualization

**heapfrag** gem, by tenderlove

Let's see what's going on in real-time

- Black pixel = unused or unallocated memory
- Green pixel = “young” object
- Red pixel = “old” object



# Memory Usage Statistics

```
[51] pry(main)> GC.stat  
=> {:count=>71,  
  :heap_allocated_pages=>4129,  
  :heap_sorted_length=>4129,  
  :heap_allocatable_pages=>0,  
  :heap_available_slots=>1682983,  
  :heap_live_slots=>1682338,  
  :heap_free_slots=>645,  
  :heap_final_slots=>0,  
  :heap_marked_slots=>1282624,  
  :heap_eden_pages=>4129,  
  :heap_tomb_pages=>0,  
  :total_allocated_pages=>4129,  
  :total_freed_pages=>0,  
  :total_allocated_objects=>10316997,  
  :total_freed_objects=>8634659,
```

```
:malloc_increase_bytes=>789296,  
:malloc_increase_bytes_limit=>25288021,  
:minor_gc_count=>57,  
:major_gc_count=>14,  
:remembered_wb_unprotected_objects=>7514,  
:remembered_wb_unprotected_objects_limit=>14268,  
:old_objects=>1266437,  
:old_objects_limit=>2336758,  
:oldmalloc_increase_bytes=>12184848,  
:oldmalloc_increase_bytes_limit=>31509677}
```

```
[50] pry(main)> GC::INTERNAL_CONSTANTS
```

```
=> {:RVALUE_SIZE=>40, :HEAP_PAGE_OBJ_LIMIT=>408,  
  :HEAP_PAGE_BITMAP_SIZE=>56,  
  :HEAP_PAGE_BITMAP_PLANES=>4}
```

# Better, Stronger, Faster

## Ruby 2.1


- Generational garbage collection

## Ruby 2.2

- Garbage collection of symbols
- Incremental garbage collection
- Support for `jemalloc`

## Ruby 2.3


- `frozen_string_literal` pragma directive






# What's Next?

- Compaction?
- Default to `jemalloc`?




# Memory Leak, or Memory Bloat?



## Memory Bloat

- Logarithmic
- “Natural” behavior of a Ruby process
- Just [download more RAM](#)
- Stepped/rapid growth is probably a code issue



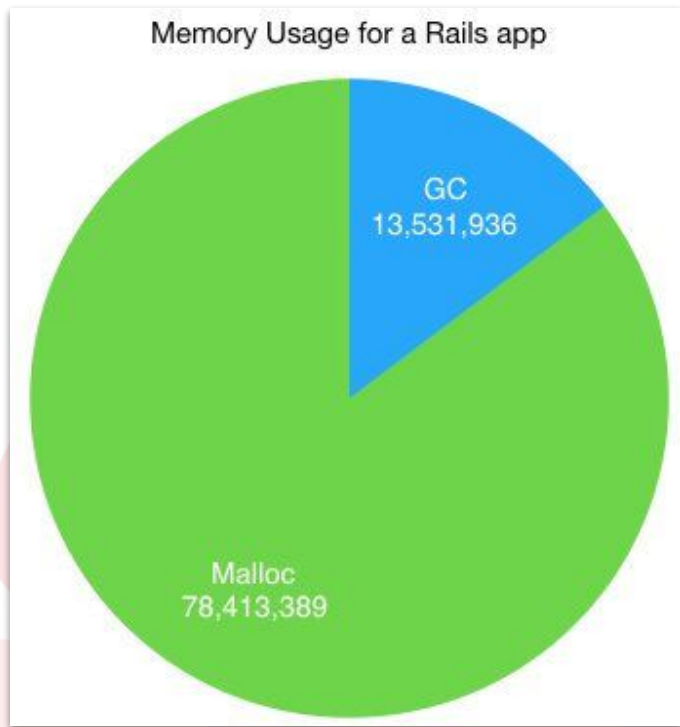
## Memory Leak

- Linear
- Will crash your application
- Far less common in Rubyland
- It's not a leak if your code is doing it intentionally

# Use Memory Efficiently

## Avoid unnecessary object allocations

- Everything is an object, thus everything your program does will allocate objects.
- This includes deallocating objects.
- Practically speaking, memory used for data outside of `ObjectSpace` will never be freed.
- Common problem: `Enumerator`
- Uncommon solution:  
`Enumerator::Lazy`



# Don't Load Data You Don't Need

- At Weedmaps, we deal with a lot of geospatial data
- Geometry data is (relatively) huge

```
[1] pry(main)> 'Texas'.length  
=> 5
```

```
[2] pry(main)> Region.find_by(name: 'Texas').geometry.to_s.length  
=> 1520683
```

- Where is that geometry data in memory?
- malloc -ed somewhere outside of ObjectSpace

# Don't Load Data You Don't Need

- Added `without_geometry` scope to avoid loading unnecessary data
- Excludes `geometry` column from `SELECT` statements

```
[1] pry(main)> Region.find_by(name: 'Texas').inspect.to_s.length  
=> 1521852
```

```
[2] pry(main)> Region.without_geometry.find_by(name:  
'Texas').inspect.to_s.length  
=> 1104
```

- Turns out it wasn't necessary for almost all requests
- Originally noticed as a query issue, not a memory one

# Tools for the Toolbox

- APM

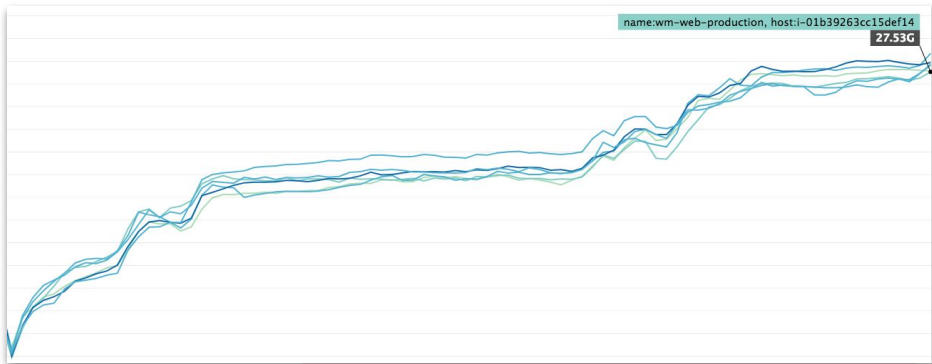
- New Relic
- Skylight
- Scout

- Gems

- `derailed_benchmarks`
- `rbtrace`
- `memory_profiler`

- You

- It's much easier to avoid memory issues in development than to troubleshoot them in a production application
- Be aware of how your code uses memory. If you're unsure, profile things before you commit them.



# Recap

- Terms and concepts
- Under the hood
  - Page allocation
  - Object allocation
  - Garbage collection
  - Visualizing the heap
- Improvements, past and future
- Troubleshooting memory consumption
  - Leak or bloat?
  - Using memory efficiently
  - Tools for the toolbox



# Questions?



@bsdzunk



[github.com/dzunk](https://github.com/dzunk)